

A Text Mining Journey through Hotel and Restaurant Reviews in Thailand.

Dataset:

The dataset contains 53,644 rows and 5 columns of customer reviews of hotels and restaurants in 25 different locations in the province of Phuket, Thailand. In this report we'll select reviews of hotels and Restaurants from two major districts in the Phuket province, each representing a customer's experience. We carefully selected 15 hotels/Restaurants located in Choeng thale in Thalang, in the Northern part of Phuket, as well as another 15 hotels/Restaurants in Cape Panwa which is situated in Mueang Phuket in the southern part of Phuket which is also the capital of the province. The dataset used for this task is the `tourist_accommodation_reviews.csv` dataset.

Explanation and preparation of dataset (Exploratory Data Analysis)

1. Open Jupyter Notebook, either from the Anaconda Prompt or the Anaconda Navigator. Open a new notebook by clicking the New tab button and select Python 3 (ipykernel) to create a new notebook workspace.
2. For this report, I will use the Natural Language Toolkit (NLTK) library, which provides a range of text-processing tools for tasks such as classification, tokenization, stemming, tagging, and sentiment analysis. Run the code below to import the libraries.

```
In [1]: #importing the necessary libraries and Packages

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import re
import seaborn as sns
from wordcloud import WordCloud
import nltk
nltk.download(['stopwords',
               'punkt',
               'wordnet',
               'omw-1.4',
               'vader_lexicon'
               ])
%matplotlib inline
```

3. Import data using the read_csv() function in the pandas library to read the dataset into a pandas data frame, and then analyse and visualise some of the results.

```
In [2]: #importing the dataset

reviews = pd.read_csv('tourist_review.csv')
```

4. To gain a better understanding of the dataset and its columns, execute the following code:
(a) reviews.head()

This function retrieves the first few rows of the dataset.

```
In [3]: #Exploration of the dataset
```

```
reviews.head()
```

```
Out[3]:
```

	ID	Review Date	Location	Hotel/Restaurant name	Review
0	m579769293	Reviewed 1 week ago	Cape Panwa	The Cove Phuket	I finally made to to The Cove after hearing fr...
1	m578446147	Reviewed 1 week ago	Cape Panwa	The Cove Phuket	The food and the service were both excellent a...
2	m578261388	Reviewed 2 weeks ago	Cape Panwa	The Cove Phuket	Almost confused with the corner bar but what a...
3	m578201696	Reviewed 2 weeks ago	Cape Panwa	The Cove Phuket	I know your probably going to reply with all t...
4	m577322860	Reviewed 2 weeks ago	Cape Panwa	The Cove Phuket	Super service and great food! Nice views and c...

(b) reviews.tail()

This function retrieves the last few rows of the dataset. To get more rows, simply specify a different number inside the parentheses. Example: “reviews.tail(20)” which would return the last 20 rows.

```
In [4]: reviews.tail()
```

```
Out[4]:
```

	ID	Review Date	Location	Hotel/Restaurant name	Review
2988	m145486095	Reviewed November 15, 2012	Choeng Thale	9' Sea Breeze	this great little place on surin beach serves ...
2989	m145450467	Reviewed November 15, 2012	Choeng Thale	9' Sea Breeze	If you're looking for a good food, great atmos...
2990	m145136705	Reviewed November 11, 2012	Choeng Thale	9' Sea Breeze	Meeting different people from all over the wor...
2991	m145067924	Reviewed November 10, 2012	Choeng Thale	9' Sea Breeze	This has to be the best sports bar restaurant ...
2992	m144513470	Reviewed November 4, 2012	Choeng Thale	9' Sea Breeze	This is a great little restaurant with loads o...

(c) reviews.describe()

This function is used to return the summary of the data.

```
In [5]: reviews.describe()
```

```
Out[5]:
```

	ID	Review Date	Location	Hotel/Restaurant name	Review
count	2993	2993	2993	2993	2993
unique	2963	1129	2	30	2963
top	m579563747	Reviewed 2 weeks ago	Cape Panwa	The Cove Phuket	We had dinner here. Just stumbled across it. S...
freq	3	32	1500	100	3

5. To identify and define stop words in the English language, run the following code:

```
In [6]: #identifying and defining stop word in english language
```

```
stop_words = nltk.corpus.stopwords.words('english')
print(stop_words)
```

```
['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'your', 'yours', 'y  
ourself', 'yourselves', 'he', 'him', 'his', 'himself', 'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',  
'they', 'them', 'their', 'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those',  
'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does', 'did', 'doing', 'a', 'an',  
'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of', 'at', 'by', 'for', 'with', 'about', 'against', 'b  
etween', 'into', 'through', 'during', 'before', 'after', 'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'of  
'f', 'over', 'under', 'again', 'further', 'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both',  
'each', 'few', 'more', 'most', 'other', 'some', 'such', 'no', 'nor', 'not', 'only', 'own', 'same', 'so', 'than', 'too', 'very',  
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're', 've', 'y', 'ain', 'ar  
en', 'aren't', 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', "hadn't", 'hasn', "hasn't", 'haven', "have  
n't", 'isn', 'isn't', 'ma', 'mightn', "mightn't", 'mustn', "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "should  
n't", 'wasn', "wasn't", 'weren', "weren't", 'won', "won't", 'wouldn', "wouldn't"]
```

6. Having explored the pre-processing steps, the texts were split into two separate parts called tokens and cleaned out by removing the stop words using the below code:

```
In [7]: #preprocessing the texts by splitting them into separate parts named tokens and cleaning out these tokens by removing the stop words

def preprocess_text(text):
    tokenized_document = nltk.tokenize.RegexpTokenizer('[a-zA-Z0-9\']+') .tokenize(text)
    cleaned_tokens = [word.lower() for word in tokenized_document if word.lower() not in stop_words]
    stemmed_text = [nltk.stem.PorterStemmer().stem(word) for word in cleaned_tokens]
    return stemmed_text
```

For Implementation in Python, We'll import the **SentimentIntensityAnalyzer** class from NLTK's **VADER**, which returns a dictionary of sentiment scores, including the overall compound score, a positive score, a negative score, and neutral scores, and retrieve the first few rows with *reviews.head()* and *reviews.tail()* for last few rows.

```
In [8]: from nltk.sentiment.vader import SentimentIntensityAnalyzer
sentiment = SentimentIntensityAnalyzer()
```

```
In [9]: reviews['compound'] = [sentiment.polarity_scores(review)['compound'] for review in reviews['Review']]
reviews['neg'] = [sentiment.polarity_scores(review)['neg'] for review in reviews['Review']]
reviews['neu'] = [sentiment.polarity_scores(review)['neu'] for review in reviews['Review']]
reviews['pos'] = [sentiment.polarity_scores(review)['pos'] for review in reviews['Review']]

reviews.head()
```

```
Out[9]:
```

	ID	Review Date	Location	Hotel/Restaurant name	Review	compound	neg	neu	pos
0	m579769293	Reviewed 1 week ago	Cape Panwa	The Cove Phuket	I finally made it to The Cove after hearing fr...	0.9388	0.000	0.685	0.315
1	m578446147	Reviewed 1 week ago	Cape Panwa	The Cove Phuket	The food and the service were both excellent a...	0.9215	0.061	0.643	0.296
2	m578261388	Reviewed 2 weeks ago	Cape Panwa	The Cove Phuket	Almost confused with the corner bar but what a...	0.8944	0.030	0.768	0.202
3	m578201696	Reviewed 2 weeks ago	Cape Panwa	The Cove Phuket	I know your probably going to reply with all t...	0.8563	0.058	0.727	0.215
4	m577322860	Reviewed 2 weeks ago	Cape Panwa	The Cove Phuket	Super service and great food! Nice views and c...	0.9549	0.000	0.479	0.521

```
In [10]: reviews.tail()
```

```
Out[10]:
```

	ID	Review Date	Location	Hotel/Restaurant name	Review	compound	neg	neu	pos
2988	m145486095	Reviewed November 15, 2012	Choeng Thale	9' Sea Breeze	this great little place on surin beach serves ...	0.8292	0.000	0.792	0.208
2989	m145450467	Reviewed November 15, 2012	Choeng Thale	9' Sea Breeze	If you're looking for a good food, great atmos...	0.9459	0.000	0.640	0.360
2990	m145136705	Reviewed November 11, 2012	Choeng Thale	9' Sea Breeze	Meeting different people from all over the wor...	0.7269	0.053	0.667	0.280
2991	m145067924	Reviewed November 10, 2012	Choeng Thale	9' Sea Breeze	This has to be the best sports bar restaurant ...	0.9796	0.000	0.572	0.428
2992	m144513470	Reviewed November 4, 2012	Choeng Thale	9' Sea Breeze	This is a great little restaurant with loads o...	0.9578	0.000	0.662	0.338

7. To gain more insights into the sentiment scores of the new columns, run the following code:

```
In [11]: #Exploring the new columns for more insights into sentiment scores for review data

reviews[['compound', 'neg', 'neu', 'pos']].describe()
```

```
Out[11]:
```

	compound	neg	neu	pos
count	2993.000000	2993.000000	2993.000000	2993.000000
mean	0.708503	0.022771	0.730721	0.246509
std	0.378836	0.045676	0.126106	0.133492
min	-0.959400	0.000000	0.286000	0.000000
25%	0.662400	0.000000	0.648000	0.147000
50%	0.865800	0.000000	0.737000	0.241000
75%	0.933600	0.037000	0.819000	0.336000
max	0.990900	0.407000	1.000000	0.714000

8. We will create a function to compute negative, neutral, and positive sentiments.

Then, add a new column named Sentiment to the dataset using the code below:

```
In [12]: #creating a funtion to compute negative, neutral and positive sentiments and add a new colum named sentiment to our dataset

def getAnalysis(score):
    if score < 0:
        return 'Negative'
    elif score == 0:
        return 'Neutral'
    else:
        return 'Positive'

reviews['sentiment'] = reviews['compound'].apply(getAnalysis)
```

9. View the dataset again to see the new columns.

```
In [13]: #viewing the new dataset
reviews.head()
```

```
Out[13]:
```

	ID	Review Date	Location	Hotel/Restaurant name	Review	compound	neg	neu	pos	sentiment
0	rn579769293	Reviewed 1 week ago	Cape Panwa	The Cove Phuket	I finally made to to The Cove after hearing fr...	0.9388	0.000	0.685	0.315	Positive
1	rn578446147	Reviewed 1 week ago	Cape Panwa	The Cove Phuket	The food and the service were both excellent a...	0.9215	0.061	0.643	0.296	Positive
2	rn578261388	Reviewed 2 weeks ago	Cape Panwa	The Cove Phuket	Almost confused with the corner bar but what a...	0.8944	0.030	0.768	0.202	Positive
3	rn578201696	Reviewed 2 weeks ago	Cape Panwa	The Cove Phuket	I know your probably going to reply with all t...	0.8563	0.058	0.727	0.215	Positive
4	rn577322860	Reviewed 2 weeks ago	Cape Panwa	The Cove Phuket	Super service and great food! Nice views and c...	0.9549	0.000	0.479	0.521	Positive

10. To view the counts for each sentiment type, run the below code.

```
In [14]: #Counts for each sentiment type

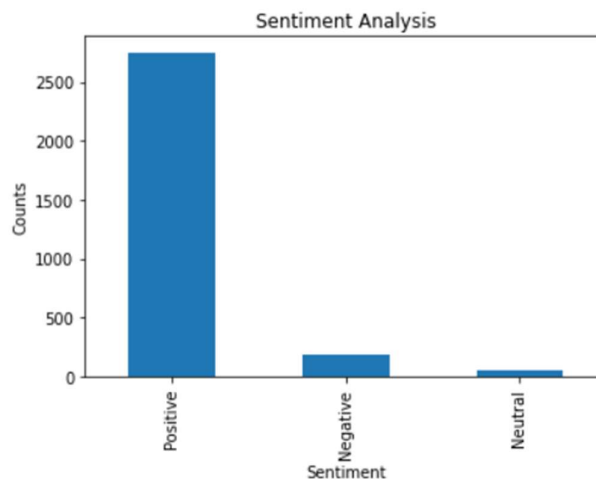
reviews['sentiment'].value_counts()
```

```
Out[14]: Positive    2749
         Negative     191
         Neutral       53
         Name: sentiment, dtype: int64
```

11. We can visualise the counts for each sentiment type in a bar graph using the following code:

```
In [15]: #visualise the counts for each sentiment type

plt.title('Sentiment Analysis')
plt.xlabel('Sentiment')
plt.ylabel('Counts')
reviews['sentiment'].value_counts().plot(kind = 'bar')
plt.show()
```




12. We can further view the distribution of compound, positive, and negative scores by executing the codes shown in the images below:

(a) Compound score

```
In [16]: #visualise distribution of compound scores  
sns.histplot(reviews['compound'])
```


```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x7f22c79bd8b0>
```



(b) Positive score

```
In [17]: #visualise distribution of positive scores  
sns.histplot(reviews['pos'])
```


```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x7f22c78cce80>
```



(c) Negative score

```
In [18]: #visualise distribution of negative scores
sns.histplot(reviews['neg'])
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f22c77ff6a0>
```



13. Let's look at how many negative reviews there are per hotel/restaurant using the following code below:

```
In [19]: #negative reviews per hotel/Restaurant
(reviews['compound']<=0).groupby(reviews['Hotel/Restaurant name']).sum()
```

```
Out[19]: Hotel/Restaurant name
360 ° Bar                    5
9' Sea Breeze                11
Ann's Kitchen Bar and Grill   5
Audy Restaurant              5
Baan Ra Tree Restaurant      12
Baba Soul Food               4
Bamboo Bar                   9
Bampot Kitchen & Bar         3
Benny's American Bar & Grill  5
Black Cat                    10
Bocconcino                   7
Bodega & Grill               9
Cafe de Bangtao              10
Chilli Kitchen               7
Curry Night Indian Restaurant 3
Cut Grill & Lounge           11
D Restaurant                 11
DaVinci Restaurant           5
DeDos                        4
Dino Park                    16
Live India Indian Restaurant  10
Mali Chic Restaurant         12
Panwa House                  8
Plum Prime Steakhouse        7
Sabai Sabai                  12
Sea Breeze                   9
The Cove Phuket              2
The Grill                    5
Tree Top Restaurant and Bar   6
Uncle Nan's Italian Restaurant 21
Name: compound, dtype: int64
```

14. Apparently, we should also look at the number of negative reviews as a proportion of the total number for each hotel/restaurant, unless it's already known. To do this, run the code in the image below:

In [20]: `#calculate the negative as a percentage of the total reviews`

```
percent_negative = pd.DataFrame((reviews['compound']<=0).groupby(reviews['Hotel/Restaurant name']).sum()  
                                /reviews['Hotel/Restaurant name'].groupby(reviews['Hotel/Restaurant name']).count()*100,  
                                columns=['% negative reviews']).sort_values(by='% negative reviews')
```

Out[20]:


Hotel/Restaurant name	% negative reviews
The Cove Phuket	2.000000
Curry Night Indian Restaurant	3.000000
Bampot Kitchen & Bar	3.000000
DeDos	4.000000
Baba Soul Food	4.000000
The Grill	5.000000
DaVinci Restaurant	5.000000
Benny's American Bar & Grill	5.000000
360 ° Bar	5.000000
Audy Restaurant	5.000000
Ann's Kitchen Bar and Grill	5.000000
Tree Top Restaurant and Bar	6.000000
Bocconcino	7.000000
Plum Prime Steakhouse	7.000000
Chilli Kitchen	7.446809
Panwa House	8.000000
Bodega & Grill	9.000000
Sea Breeze	9.000000
Bamboo Bar	9.000000
Live India Indian Restaurant	10.000000
Black Cat	10.000000
Cafe de Bangtao	10.101010
D Restaurant	11.000000
Cut Grill & Lounge	11.000000
9° Sea Breeze	11.000000
Mali Chic Restaurant	12.000000
Sabai Sabai	12.000000
Baan Ra Tree Restaurant	12.000000
Dino Park	16.000000
Uncle Nan's Italian Restaurant	21.000000

15. We can further plot the above percentage as a horizontal barplot using seaborn.

In [21]: `#we can also plot the above percentages as a horizontal barplot`

```
sns.barplot(data=percent_negative, x='% negative reviews', y=percent_negative.index, color='c')
```

Out[21]: `<matplotlib.axes._subplots.AxesSubplot at 0x7f22c76ceaf0>`



We can see that some of the words aren't particularly informative. for example: 'food', 'pizza', 'Italian, and 'pasta' are all references to the restaurant menu. But we can also see words like 'disappoint', 'worst', 'expensive', and 'taste' have been mentioned a few times – Maybe these are issues that require further investigation.

18. We can also generate a wordcloud from the positive reviews using the code below.

```
In [24]: #wordcloud of words from positive reviews for Uncle Nan's Italian Restaurant

pos_tokens = [word for review in reviews_positive_subset['processed_review'] for word in review]

wordcloud = WordCloud(background_color='white').generate_from_text(
    ' '.join(pos_tokens))

plt.figure(figsize=(12,12))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")
plt.show()
```



19. Wordclouds do provide a way to visualise word frequencies, but sometimes it might be difficult to explain. we can also use the tabulate method to understand the most frequent words, and the number of occurrences in each with the use of FreqDist from NLTK using the code below:

(a) Positive Reviews

```
In [25]: #using FreqDist from NLTK to show the frequency of words in a tabular form for positive reviews

from nltk.probability import FreqDist

pos_freqdist = FreqDist(pos_tokens)

pos_freqdist.tabulate(10)
```

food	good	italian	pizza	servic	nice	staff	restaur	meal	hotel
59	34	31	26	21	19	18	17	17	16

(b) Negative Reviews

In [26]: *#using FreqDist from NLTK to show the frequency of words in a tabular form for negative reviews*

```
from nltk.probability import FreqDist
neg_freqdist = FreqDist(neg_tokens)
neg_freqdist.tabulate(10)
```


food	italian	pizza	restaur	servic	order	went	pasta	disappoint	slow
18	10	9	7	7	6	6	6	5	5

20. We can also use the frequency distribution graph to show word frequency in both positive and negative reviews using the following codes below:

(a) Positive reviews

In [27]: *#Frequency distribution graph to show word frequency for postive reviews*

```
pos_freqdist.plot(30)
```




Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x7f22c7fb9fd0>

(b) Negative reviews

In [28]: *#Frequency distribution graph to show word frequency for negative reviews*

```
neg_freqdist.plot(30)
```



Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x7f22c795cee0>

Results analysis and discussion

Our dataset did not require any pre-processing as we used the implementation of VADER, which takes care of the process. We utilised tokenization, stopword removal, and stemming to generate a wordcloud that visualises the distribution of words and word count. This helped to identify the main words that contribute to positive and negative reviews at Uncle Nan's Italian Restaurant. The management can use this information to investigate and improve specific areas.

Conclusions

For the implementation of this analysis, we chose VADER due to its efficiency, ease of use, and clear, compelling results. This report aims to assist restaurants and hotels in Thailand in identifying their strengths and weaknesses, allowing them to improve their quality of service in the right areas and ultimately leading to higher positive scores.